

Humans reciprocate intentional acts by discriminating against group peers^{*}

David Hugh-Jones[†] Itay Ron[‡] Ro'i Zultan[§]

This version: July 30, 2016.

Abstract

Cycles of intergroup revenge appear in large scale conflicts. We experimentally test the hypothesis that humans practice group-based reciprocity: if someone harms or helps them, they harm or help other members of that person's group. Subjects played a trust game, then allocated money between other people. Senders whose partners returned more in the trust game gave more to that partner's group members. The effect was about 60 per cent of the size of the direct reciprocity effect. Receivers' allocations to group members did not depend on their partner's play, suggesting that group reciprocity was only triggered when the partner's intentions were unequivocal.

1 Introduction

The evolution of cooperation poses a puzzle to evolutionary and social scientists. Cooperation—by which individuals forgo personal benefits to aid others—is a hallmark of human uniqueness. However, there is a direct evolutionary pressure selecting against cooperators in favour of free riders, who can reap the societal benefits of cooperation without paying the costs. One mechanism that may underlie the

^{*}This research was supported by the ISRAEL SCIENCE FOUNDATION (grant No. 214/13)

[†]School of Economics, University of East Anglia. E-mail: d.hugh-jones@uea.ac.uk.

[‡]E-mail: itayron@gmail.com

[§]Department of Economics, Ben-Gurion University. E-mail: zultan@bgu.ac.il.

evolution of cooperation is reciprocity (**nowak2006five**; **nowak2012evolving**). Direct reciprocity increases the fitness of cooperators, as helping another individual leads to that individual reciprocating the help. In indirect reciprocity, individuals help or harm other people than those who have helped them.

Indirect reciprocity comes in two flavours: *downstream* reciprocity follows the maxim ‘do unto thy neighbour as they have done to others’, whereas *upstream* reciprocity follows the maxim ‘do unto thy neighbour as others have done unto you’. In this paper we study upstream indirect reciprocity (**boyd1989evolution**; **nowak2005evolution**; **nowak2007upstream**). Compared to downstream reciprocity, upstream reciprocity is simpler to implement, as it does not require tracking individual reputations, but is more difficult to understand from an evolutionary point of view (**boyd1989evolution**; **nowak2005evolution**). Nonetheless, upstream reciprocity can co-evolve with direct or spatial reciprocity (**nowak2007upstream**). Furthermore, laboratory experiments provide positive evidence for upstream reciprocity, as individuals are more generous to others if a third party was generous to them (**dufwenberg2001direct**; **guth2001trust**; **greiner2005indirect**), and the mere possibility of being harmed by a third party reduces cooperation in a social dilemma (**weisel2016social**).

Direct and indirect reciprocity are embedded in a socially structured world. Evidence suggests that human social behaviour—including cooperation and reciprocity—is bounded by social groups (**tajfel1979integrative**; **yamagishi2000thegroup**; **balliet2014ingroup**; **DeDreu2014**). Humans cooperate more with, and reciprocate more towards members of their group (**chen2009group**; **chen2011potential**); and they punish norm violators more if these harm a member of their own group (**bernhard2006group**; **bernhard2006parochial**).

Here we examine group-based upstream reciprocity, or *group reciprocity*. That is, an individual who is harmed (helped) by a member of an outgroup becomes more likely to harm (help) others from that group. Whereas group-based downstream reciprocity (**bernhard2006group**; **bernhard2006parochial**) follows the maxim ‘do unto others as they have done to members of *my* tribe’, group-based upstream reciprocity follows the maxim ‘do unto others as members of *their* tribe have done to me’ (Figure ??). Both kinds of motive can expand the scope of an individual-level conflict into an inter-group conflict. While (group-based) downstream reci-




Figure 1: Upstream reciprocity. (a) Someone who was helped or harmed becomes equally more likely to help or harm new partners. (b) Upstream reciprocity is particularly directed at new partners that belong to the same group as the initial partner.

procuity can bring a victim’s groupmates in as retaliators, upstream reciprocity makes members of the aggressor’s group into potential targets. Ethnic conflicts provide ample evidence that group identities play an important role in generalized reciprocity. Humans take revenge on outgroups as well as on individuals, as evidenced by the cycles of intergroup revenge that appear widespread in human conflict ([horowitz2001thedeaddy](#); [horowitz1985ethnicgroups](#); [chagnon1988lifehistories](#); [shayo2010judicial](#)).

Upstream group reciprocity also has different cognitive requirements from related phenomena. While in-group altruism and group-based downstream reciprocity require people to differentiate their own group from outsiders—“us” from “them”—upstream group reciprocity requires them to differentiate between different outgroups—between “them and them”—and to keep a mental account of outgroups’ reputation. Upstream group reciprocity could thus provide an evolutionary basis for the phenomenon of outgroup stereotyping .

We ran a laboratory experiment to test the hypothesis that people reciprocate towards groups. After an initial group-formation stage, participants interacted in two strategic stages. The first stage represented the upstream interaction, in which the individual could be helped by another person. We implemented this as a Trust Game (TG) ([berg1995trust](#)), in which the Sender (S) receives 150 money-equivalent tokens, and chooses how many of them to send to the Responder (R).

The amount sent is multiplied by a factor of 3, so that R receives between 0 and 450 tokens, of which he can send any number back to S. The TG enables us to model two types of interactions. Whereas R is clearly kind when returning money (and nasty when exploiting a generous proposer by keeping the received amount), S's intentions are equivocal. Sending money can be driven by selfish expectations of reciprocity, while not sending can be driven by caution. Thus, R experiences kind or nasty *actions*, while S experiences kind or nasty actions and *intentions*.

The second stage represented the reciprocal action, in which the individual could help others. We implemented this interaction as an allocation game, in which an allocator allocates a fixed amount between two recipients. In Direct Reciprocity rounds, the recipients included the TG partner; in Group Reciprocity rounds, a member of the TG partner's group; and in In-Group Favouritism rounds, a member of the allocator's group. The other recipient was always a member of a third, neutral, group. Baseline rounds included two neutral recipients, to test whether the TG experience lead allocators to discriminate in absence of any reciprocal or group motivations. This design is an advance on previous work ([gaertner2008whenrejection](#); [stenstrom2008theroles](#); [hugh-jones2013intergroup](#)), because it tests for group reciprocity using material incentives, clearly differentiates the individual partner from his or her group members, and rules out strategic motives.

2 Method

Each session consisted of 24 participants, who were randomly allocated into six groups of four. To distinguish from the groups in which participants interacted throughout the different stages of the experiment, we refer to the identity-relevant initial group as a *Team*. Each participant was identified throughout the experiment by team colour and individual number (1–4) within the team. At the beginning of the experiment, we informed participants that the experiment has five distinct stages. In each stage, new groups are formed, and it is possible (though not necessary) that two participants will be in the same group in two different stages. The specific instructions for each stage were distributed and read aloud at the beginning of the stage. The five stage were a group formation stage, the TG stage,

the AG stage, a social value orientation elicitation (**murphy2011measuring**) stage and a collectivism scale measurement (**Singelis1995horizontal**).

Following **chen2009group**, we created group identity in the first stage by allowing participants to consult each other using anonymous chat while solving a simple task. Participants solved five Raven matrices (see supplementary material). Each matrix was presented on screen for 120 seconds, during which each participant could both send written messages to the team and update her own answer. The final answer submitted at the end of the 120 seconds determined payoffs, with 10 tokens paid for each correct answer. To further boost group identity through a common goal, team members each earned an additional bonus of 5 tokens if all four team members answered correctly.

Next, participants were rematched into pairs to play the one-shot TG. To facilitate understanding, participants played five practice rounds, in which they entered decisions both as S and as R. In the actual interaction, participants could see their TG partner's team colour and individual number. The kindness of the Sender was measured as the share of the endowment sent, and of the Responder as the share of the received amount sent back. Responder's kindness was not defined for six (out of 96) responders whose partner did not send any money.

The third stage consisted of six rounds. In each round, participants interacted in a group of three, identified by team colour and number. The allocation game was implemented as a random dictator game as follows. Each player in the group of three received 100 tokens to allocate within the group. The allocator received 30 tokens, and could freely allocate the remaining 70 tokens between the other two players. Previous research found that people do not harm, but refrain from helping negatively perceived out-groups (**weisel2015ingroup**). Accordingly, we set the parameters of the game so that, compared to the reference point of the allocator's own share, an equal division benefits both other players. The matching scheme ensured that over the six rounds, each participant was in the same group of three exactly once with a member of her own team (treatment IG), once with her TG partner (Treatment DR), and twice with other members of the TG partner's team (treatment GR). The remaining two rounds served as the baseline (B) treatment. Note that the matching is not independent. For example, if one player is in treatment DR, then the other two are in treatments DR and either B or GR. No

feedback was provided between rounds. The payoffs for this stage were determined by one randomly chosen round of the six rounds, and the allocation decision of one randomly chosen player in each group.

The fourth stage implemented the slider measure of social value orientation (**murphy2011measuring; crosetto2012flexible**), in which participants choose nine allocations between themselves and another person. For consistency with the previous stages, the team identity of the partner was known. To keep the decision independent of previous experience with the different teams, we matched participants within teams. Consistently with the notion of humans as parochial altruists, we believe that such in-group allocations best reflect social preferences. Payoffs were determined by one randomly chosen decision of the nine decisions made by one randomly chosen player in each dyad. The decisions yielded a social orientation angle for each participant, with 0° corresponding to selfishness, 45° to pure altruism, and negative angles to spitefulness.

After the fifth and final stage (a non-strategic and non-incentivised collectivism measurement), participants learned their cumulative payoff in tokens, and final payment in New Israeli Shekels. One hundred and ninety two participants, recruited using ORSEE (**greiner2015subject**) participated in eight sessions conducted between June 2014 and January 2015. Due to software malfunction in one session, the data from the AG sixth round were not saved. The experiment was programmed in z-Tree (**Fischbacher2007**). The average payment was 72 NIS (approximately \$18) for a duration of 70 minutes.

3 Results

We report results on allocations, discrimination between recipients, and direct and group reciprocity. All reported statistical tests are based on mixed-effects regressions with robust standard errors clustered on subjects. See the supplementary material for the full specification and results.

The second column in Table ?? presents the mean allocations. Participants gave significantly more to members of their own team at the expense of the neutral recipient ($z = 5.55, p < 0.005$ for senders, $z = 4.79, p < 0.001$ for responders), establishing that our group formation manipulation was successful in inducing

Table 1: Allocations and discrimination

Senders			
	Discrimination	Mean allocation	Reciprocity
Baseline	4.06 (0.99)	35.00 —	0.00 —
Group	8.10*** (1.73)	34.58* (0.78)	8.76** (4.34)
Direct	22.13*** (2.84)	33.59* (1.82)	13.89* (7.95)
In-Group	15.72*** (2.42)	39.25** (1.34)	0.71 (5.63)
Observations	564	728	728
Responders			
	Discrimination	Mean allocation	Reciprocity
Baseline	2.14 (0.64)	35.00 —	0.00 —
Group	6.18*** (1.32)	34.59 (0.68)	1.56 (1.82)
Direct	22.43*** (2.83)	35.47 (1.63)	22.76*** (4.59)
In-Group	17.31*** (2.46)	41.76*** (1.41)	1.43 (3.66)
Observations	564	752	752

Notes: Mean discrimination, mean allocation, and reciprocity (marginal effect of TG partner's kindness on allocations) by recipient's identity. Estimates are based on mixed-effects regressions with robust standard errors clustered on subjects. Separate regressions for discrimination, Senders' allocations, and Responders' allocations. Discrimination is higher than between two neutral players in all three experimental conditions. Allocations to in-group members are consistently higher than for a neutral player. Allocations to the TG partner strongly depend on his TG behavior. Allocations to the TG partner's team mates is dependent on his TG behavior only for Senders. Significance of comparison to Baseline is marked. *, **, and *** indicate $p < 0.10$, $p < 0.05$, and $p < 0.01$, respectively.

group identity and triggering in-group favouritism. Allocations to the TG partner and his team mates were not significantly different to the baseline 35 ($p > 0.440$ for all comparisons). Nonetheless, allocators discriminated significantly more than in the baseline both when interacting with their TG partner ($z = 9.36, p < 0.001$) and with his team mates ($z = 4.56, p < 0.001$), as reported in the first column of Table ???. The interaction of the identity of the recipient and the allocator’s TG role (sender or receiver) was not significant ($\chi^2 = 1.42, p = 0.700$).

Direct and group reciprocity

Figure ?? presents the allocations made in the DR and GR conditions separated by the whether the kindness of the action of the TG partner was below the median or not. The median kindness of the senders was $\frac{1}{2}$ (median transfer of 75 out of 150 tokens). The median kindness of the responders was $\frac{1}{3}$ (median share returned matches the amount sent). The figure reveals strong direct reciprocity. Senders allocate less than the baseline 35 tokens to responders who failed to at least compensate the sender for the amount sent. Responders both allocate more than 35 tokens to senders who sent at least half their endowment, and less to those who sent less. These observations are confirmed by the regression analysis (see third column in Table ??). Allocations to the TG partners increase with that partner’s kindness both for senders ($z = 1.75, p = 0.080$) and for responders ($z = 4.96, p < 0.001$).

Group reciprocity, however, is only observed for senders, who allocate less to team mates of a responder who failed to at least compensate the sender for the amount sent—a clear signal of bad intentions. Responders, on the other hand, although directly reciprocating the TG partner, do not systematically discriminate against team mates of a sender who sent little—an unkind action that does not unequivocally signals a bad intention. The regression analysis shows no significant effect of sender kindness on responder’s allocation to the sender’s team mates ($z = 0.86, p = 0.391$). The responder’s kindness, on the other hand, significantly increases allocations made to his team mates ($z = 2.02, p < 0.05$). The estimated ratio of the group and direct reciprocity coefficients is 0.63, implying that for every allocation dollar that the responder loses due to an unkind action in the TG, his

team mates lose 63 cents.

4 Discussion

Our results show that upstream reciprocity is moderated by social boundaries. We distinguish between two types of reciprocity: reciprocated actions and reciprocated intentions (**stanca2009testing**). People discriminate against others who harm them even if the harmful action does not necessarily indicate bad intentions. However, they only generalize to the perpetrator’s group members if the intentions behind the harmful actions are unequivocally bad.

This observation raises new questions regarding the nature of reciprocity and the role of intentions (or perceptions thereof). One possible interpretation stems from the distinction between intention-based and outcome-based motives in reciprocal behaviour (**falk2006theory**). Whereas intention-based reciprocity is plausibly generalisable, i.e. subjects may wish to punish there is no reason to generalize outcome-based reciprocity to group members who don’t share the outcome. The notion of outcome-based reciprocity is also related to that of restorative justice. A desire to equalize earnings given the TG behaviour in the experiment can motivate discrimination against a misbehaving TG partner, but not his fellow group members.

A related distinction is between actions and types. Stereotypical thinking can easily lead people to generalise types, whereas generalisation of actions is nonsensical. The conjecture ‘One member of the Blue group is a bad person, therefore all Blue members are bad’ is plausible. The conjecture ‘One member of the Blue group did not send any money, therefore other Blue members did not send money’ is not—as the other Blue members was not even necessarily asked to send any money.

Upstream reciprocity is notoriously difficult to understand in evolutionary terms (**boyd1989evolution**; **nowak2007upstream**) Group reciprocity may provide another piece of the puzzle, as it provides two new channels by which upstream reciprocity may evolve. First, group members are interdependent, especially in the small groups that were the norm during most of human evolutionary history. Punishing a perpetrator’s group member therefore indirectly harms the perpetrator,




Figure 2: Allocations made by to the TG partner (Direct Reciprocity) and his team mates (Group Reciprocity) by TG partner's kindness with 95% confidence intervals based on subject-level averages. Responders give more (less) to senders who sent more (less) than the median (equal to exactly half the endowment)—but do not generalize to the sender's team mates. Senders give less to responders who returned less than the median (equal to exactly $\frac{1}{3}$, i.e., the amount sent) and their team mates.

who is dependent on his peers for, e.g., public goods provision. Thus, group reciprocity may bridge upstream indirect reciprocity and direct reciprocity.

Second, the evolution of indirect reciprocity acts by way of chains of reciprocal actions, which return with some probability to the original instigator of the chain (**nowak2007upstream**). In a population organised in groups, such that people interact more frequently with their own group members, group reciprocity may increase the likelihood of successful reciprocal chains, facilitating the evolution of upstream reciprocity. We aim to explore these ideas in future work.